

Real-time Dynamic Function Call Detection

A heuristic-based approach

Franck de Goër

UGA, VU Amsterdam
ANSSI

Sanjay Rawat

VU Amsterdam

Dennis Andriesse

VU Amsterdam

Herbert Bos

VU Amsterdam

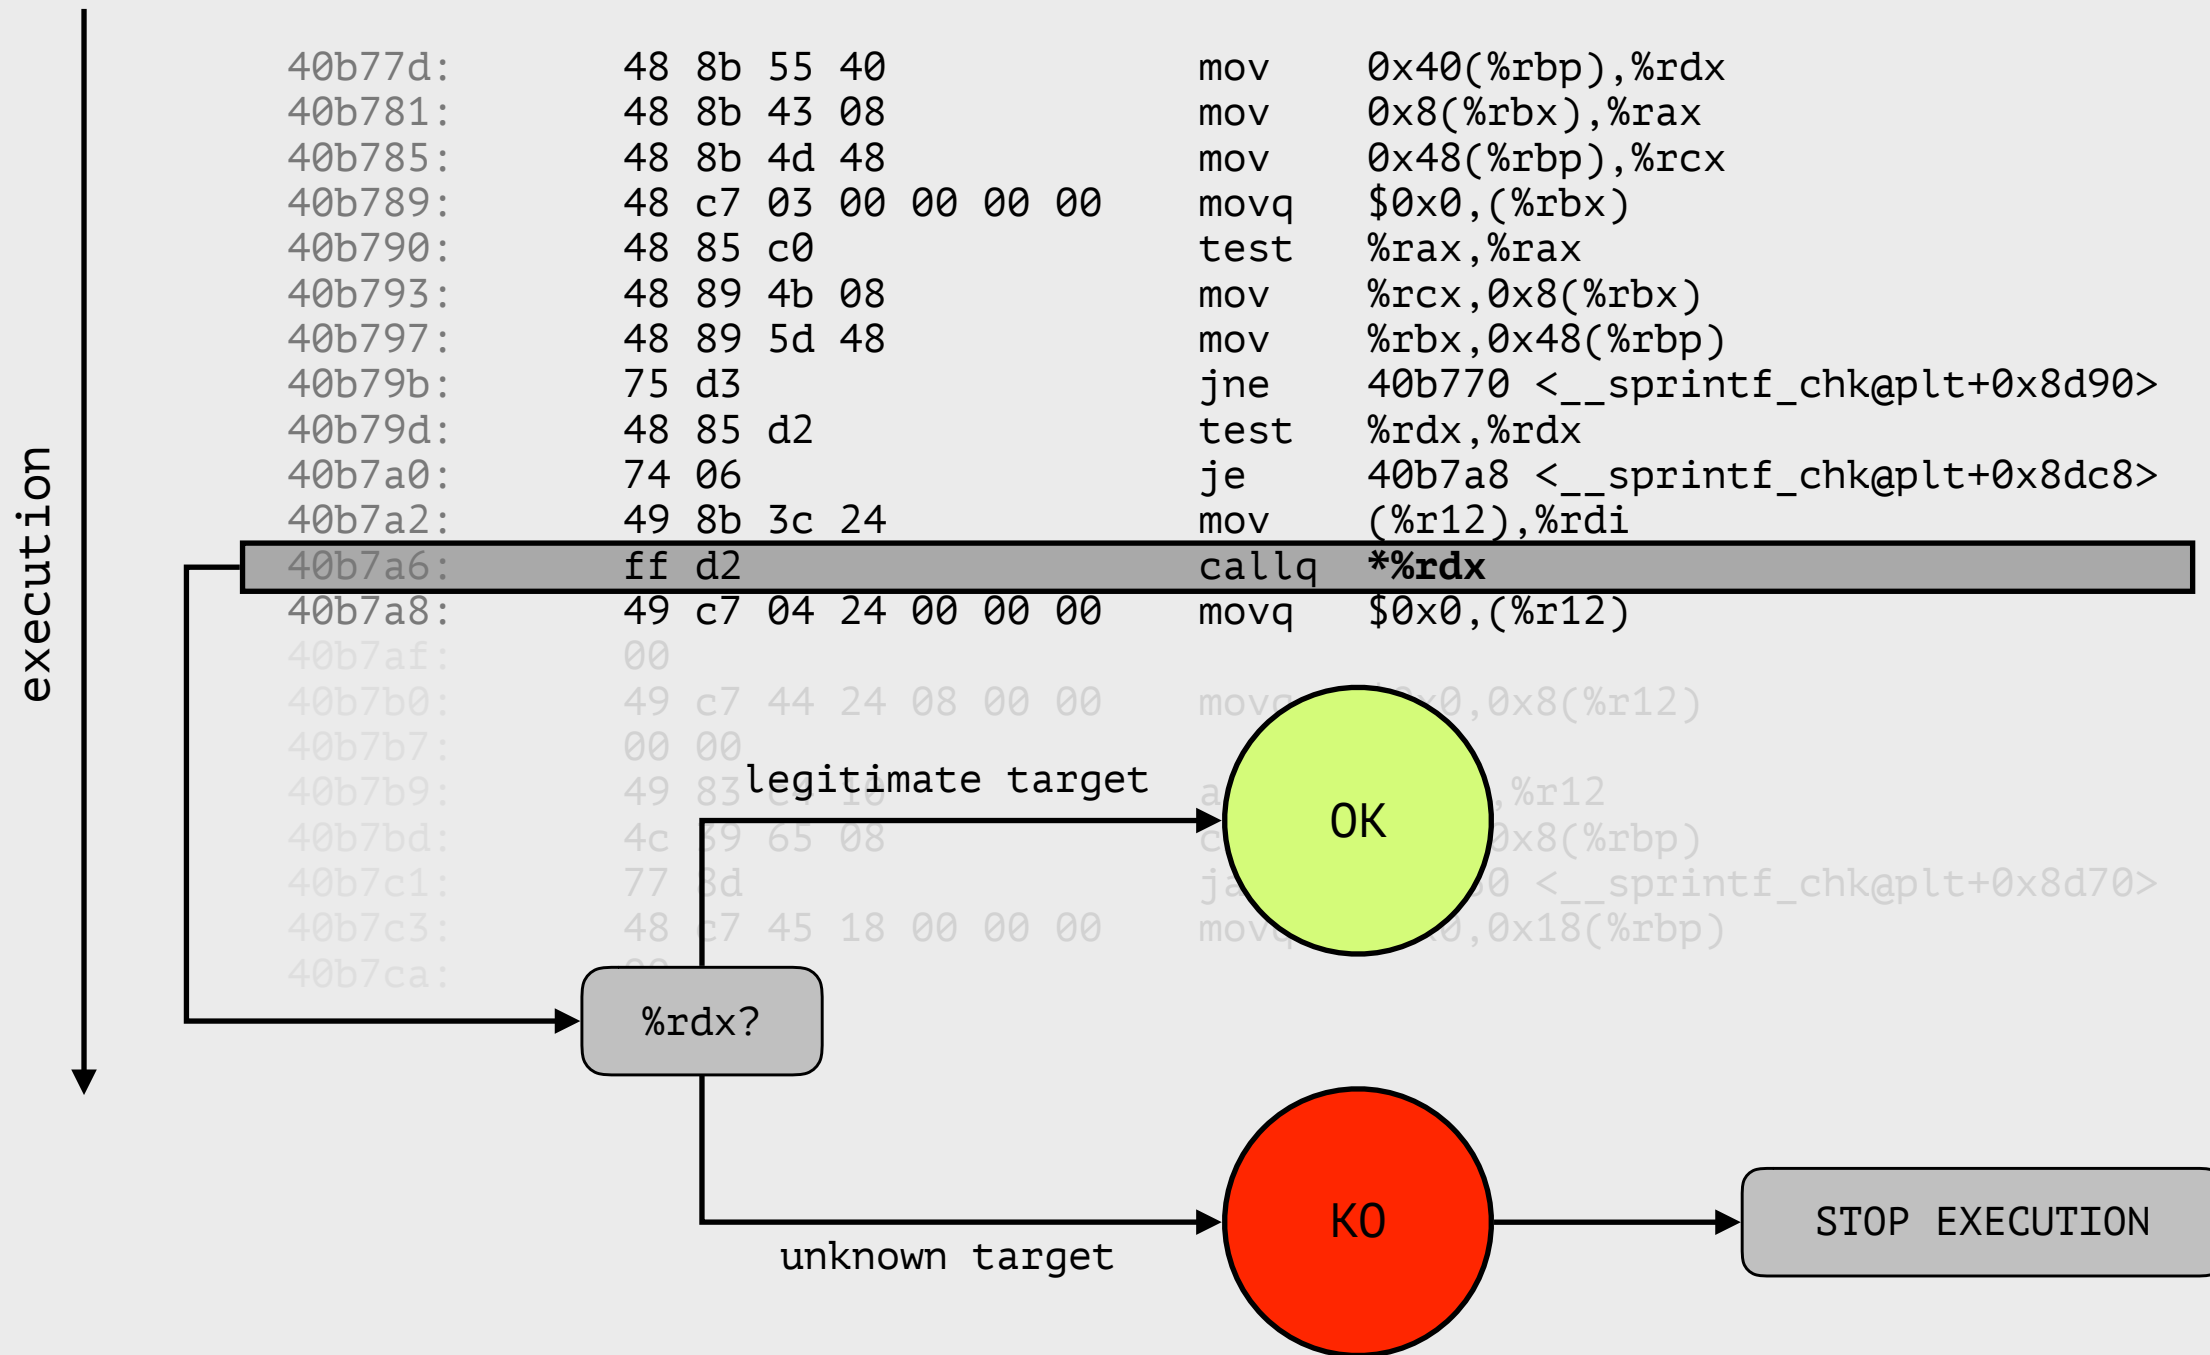
Roland Groz

UGA

INTRODUCTION

let's try to keep everyone awake
for two minutes

AN EXAMPLE TO ILLUSTRATE THE PROBLEM – CFI



AN EXAMPLE TO ILLUSTRATE THE PROBLEM – CFI

execution ↓	40b77d:	48 8b 55 40	mov	0x40(%rbp),%rdx
	40b781:	48 8b 43 08	mov	0x8(%rbx),%rax
	40b785:	48 8b 4d 48	mov	0x48(%rbp),%rcx
	40b789:	48 c7 03 00 00 00 00	movq	\$0x0,(%rbx)
	40b790:	48 85 c0	test	%rax,%rax
	40b793:	48 89 4b 08	mov	%rcx,0x8(%rbx)
	40b797:	48 89 5d 48	mov	%rbx,0x48(%rbp)
	40b79b:	75 d3	jne	40b770 <__sprintf_chk@plt+0x8d90>
	40b79d:	48 85 d2	test	%rdx,%rdx
	40b7a0:	74 06	je	40b7a8 <__sprintf_chk@plt+0x8dc8>
	40b7a2:	49 8b 3c 24	mov	(%r12),%rdi
			push	40b7a8
			jmp	*%rdx
	40b7a6:	ff d2	callq	*%rdx
	40b7a8:	49 c7 04 24 00 00 00	movq	\$0x0,(%r12)
	40b7af:	00		
	40b7b0:	49 c7 44 24 08 00 00	movq	\$0x0,0x8(%r12)
	40b7b7:	00 00		
	40b7b9:	49 83 c4 10	add	\$0x10,%r12
	40b7bd:	4c 39 65 08	cmp	%r12,0x8(%rbp)
	40b7c1:	77 8d	ja	40b750 <__sprintf_chk@plt+0x8d70>
	40b7c3:	48 c7 45 18 00 00 00	movq	\$0x0,0x18(%rbp)
	40b7ca:	00		

CFI **SHOULD** APPLY, BUT WILL IT?

THE PROBLEM

HOW TO INSTRUMENT EVERY CALL
DURING EXECUTION?

THE REAL PROBLEM

HOW TO DISTINGUISH INTRA-PROCEDURAL
JUMPS FROM JUMP-BASED CALLS?

THE REAL REAL PROBLEM

HOW TO DISTINGUISH INTRA-PROCEDURAL
JUMPS FROM TAIL CALLS?

OUR PROPOSED SOLUTION: iCi

	-00	-01	-02	-03
	iCi	iCi	iCi	iCi
binutils	1,000	1,000	1,000	1,000
evince	0,991	0,985	0,985	0,986
coreutils	0,999	0,998	0,998	0,998
ffmpeg	0,997	0,997	0,997	0,997
SPEC2006 jcall worst	nc	nc	1,000	nc
SPEC2006 iCi worst	nc	nc	0,947	nc

fscore

APPROACH

approach

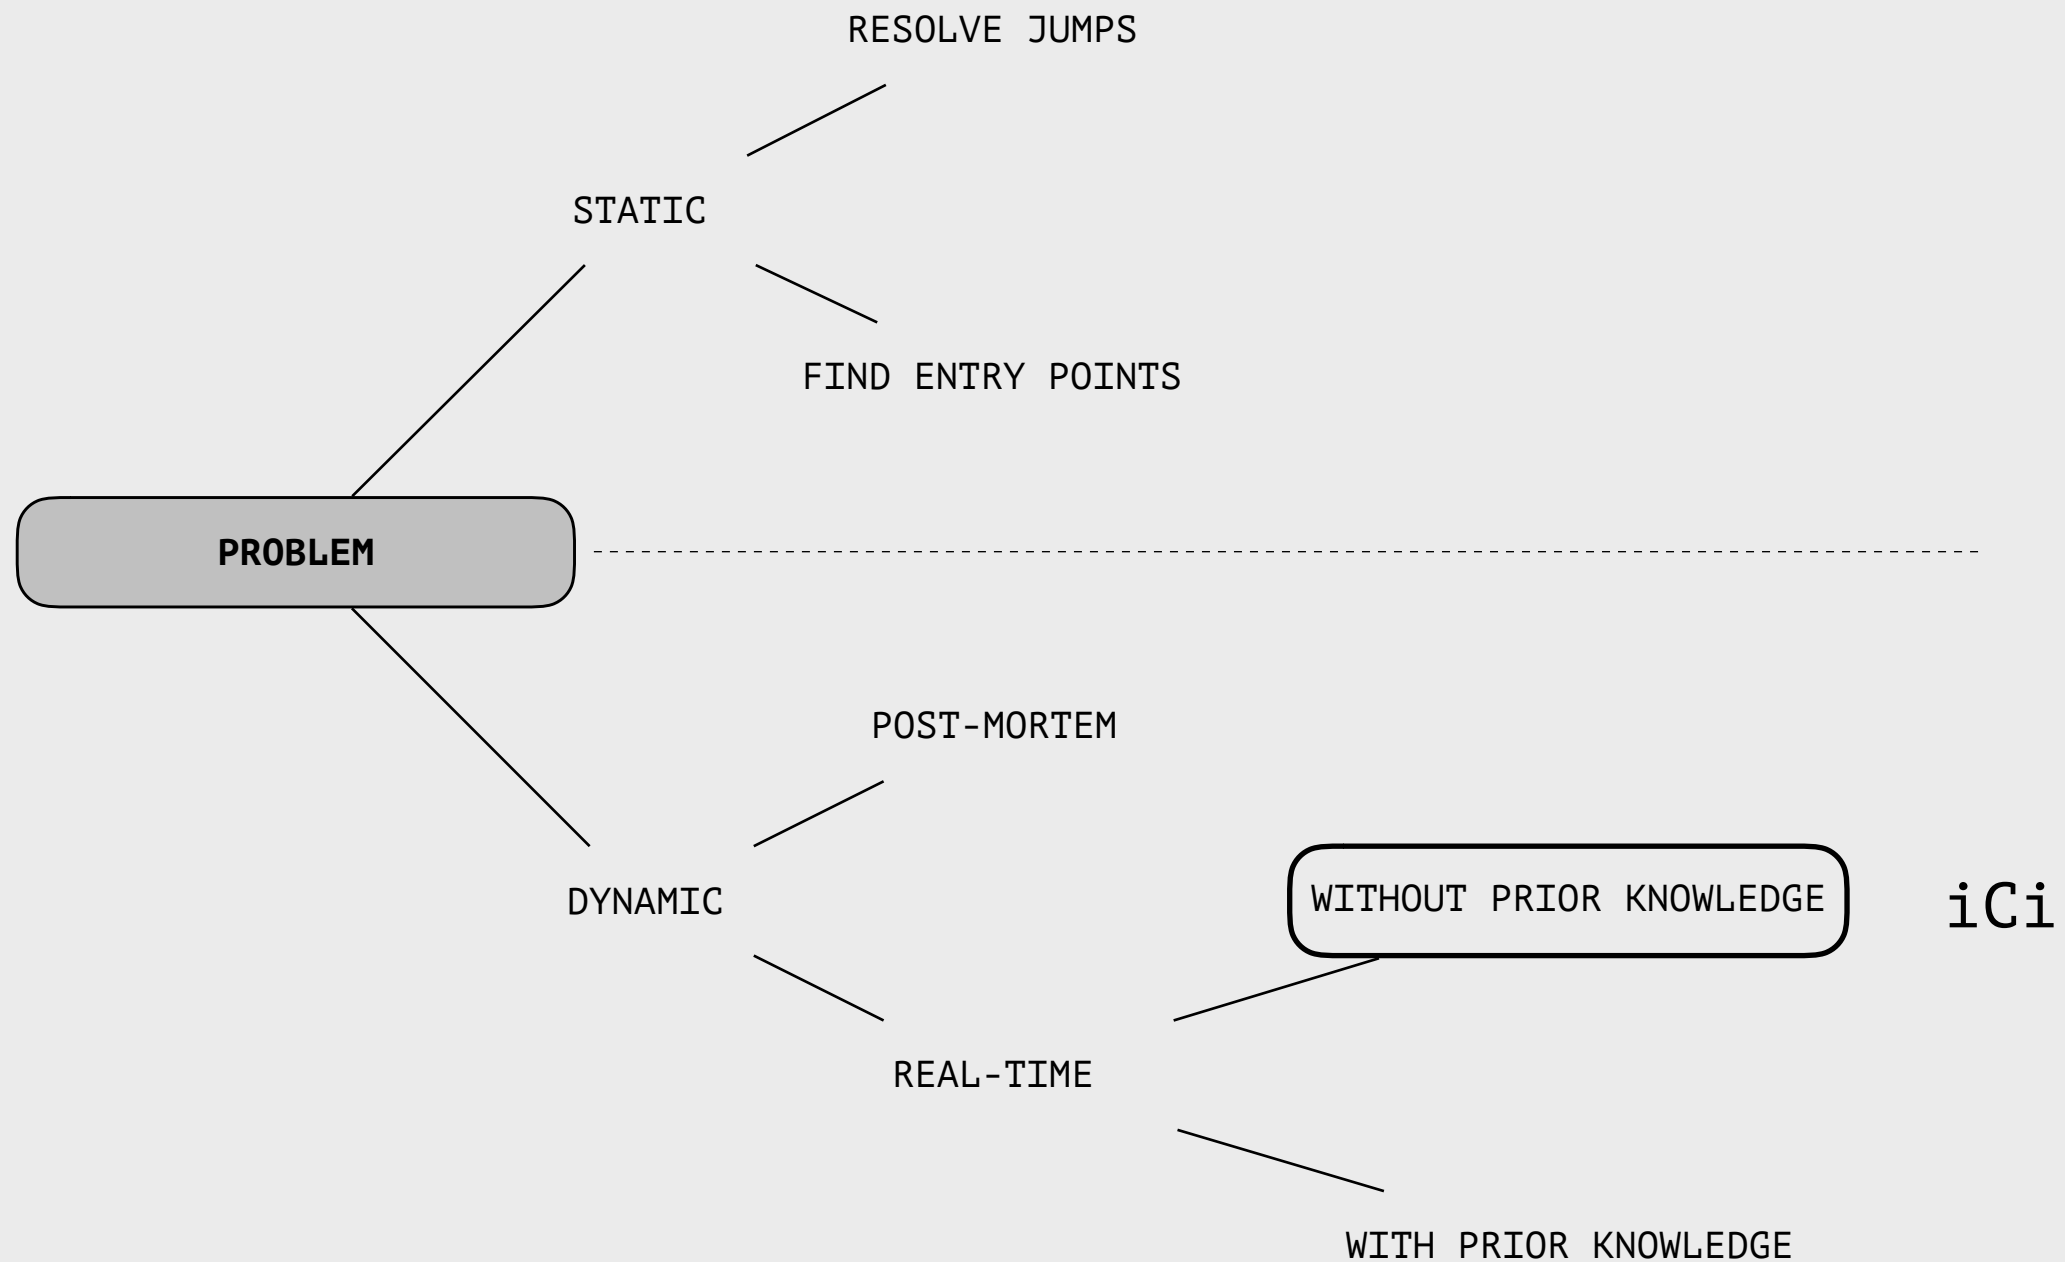
heuristics

experiments

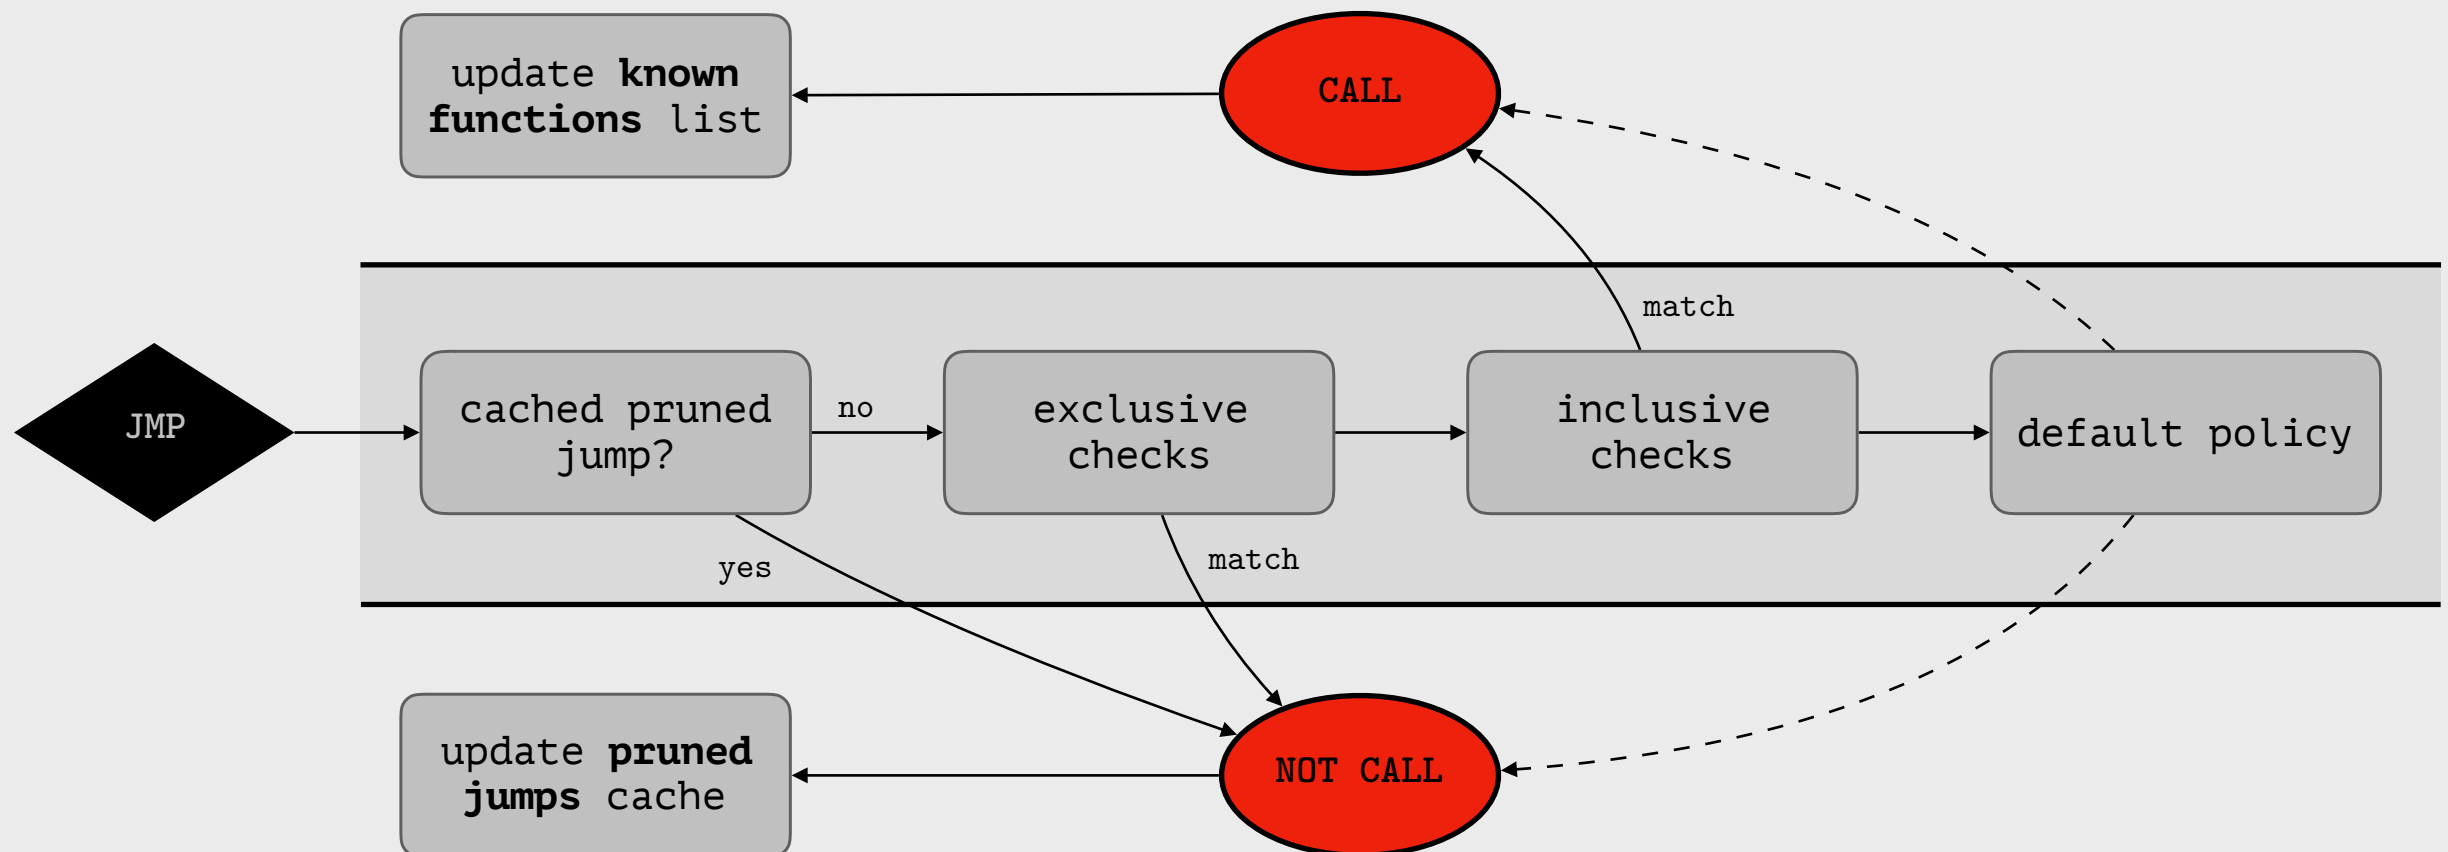
overview

OVERVIEW

POSSIBLE APPROACHES



JMP SELECTION IN iCi



approach

heuristics

experiments

overview

BACKGROUND INFORMATION

LIST OF ENTRY POINTS

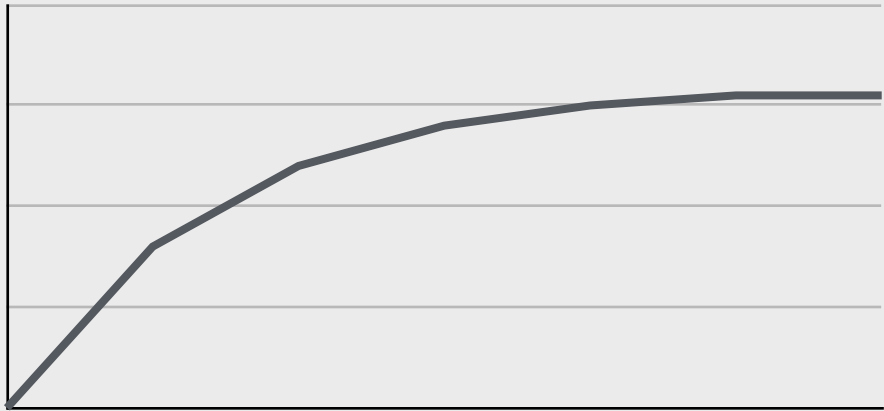
Each time a function call is detected, its **target** is marked as an **entry point**.

LIST OF RETURN LOCATIONS

Each time a function returns, the location of the last instruction is marked as a **return point**.

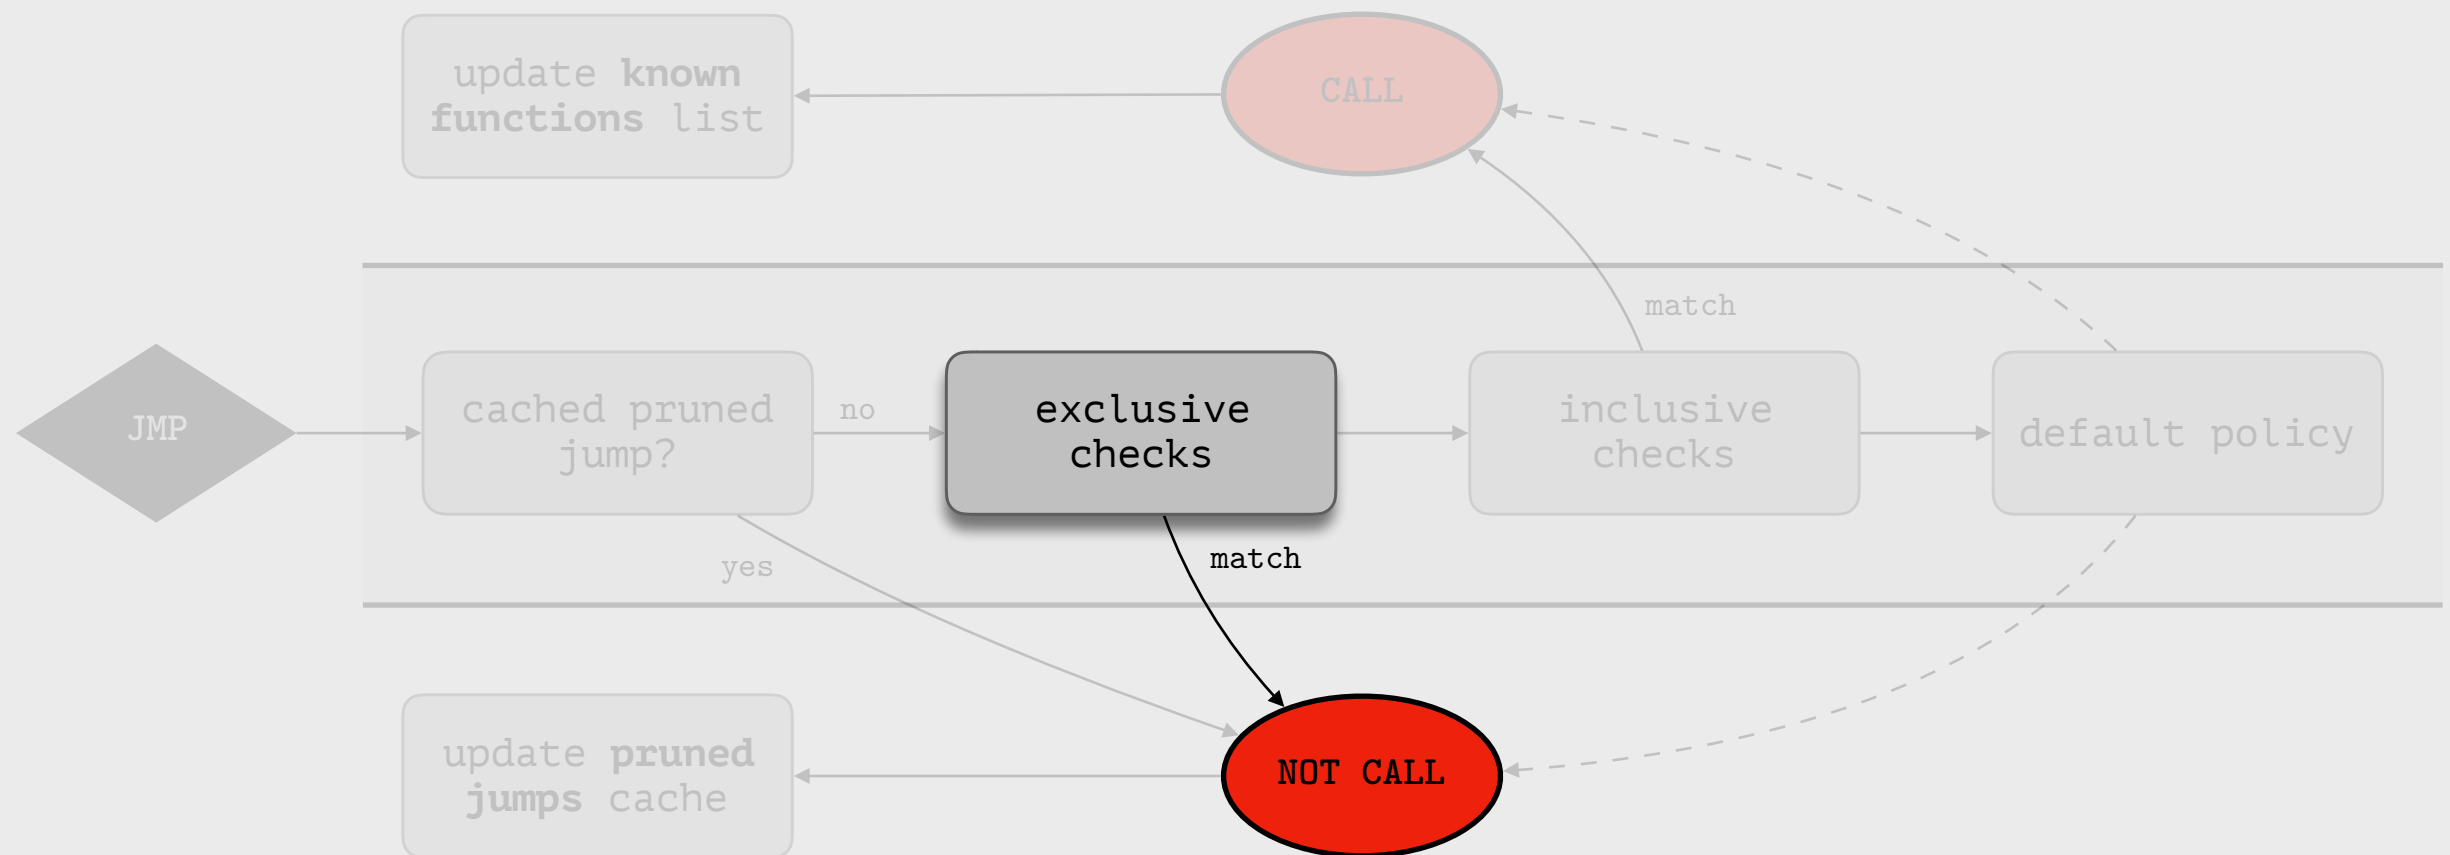
AVAILABLE INFO FOR EACH JMP

information	accessible?
target of the JMP	ok (dynamic)
current %rip, %rsp	ok (dynamic)
current entry point	~ok (depend on the known entry points)
state of stack when last call occurred	~ok (idem)
return locations for the current function	not at first call; then ~ok



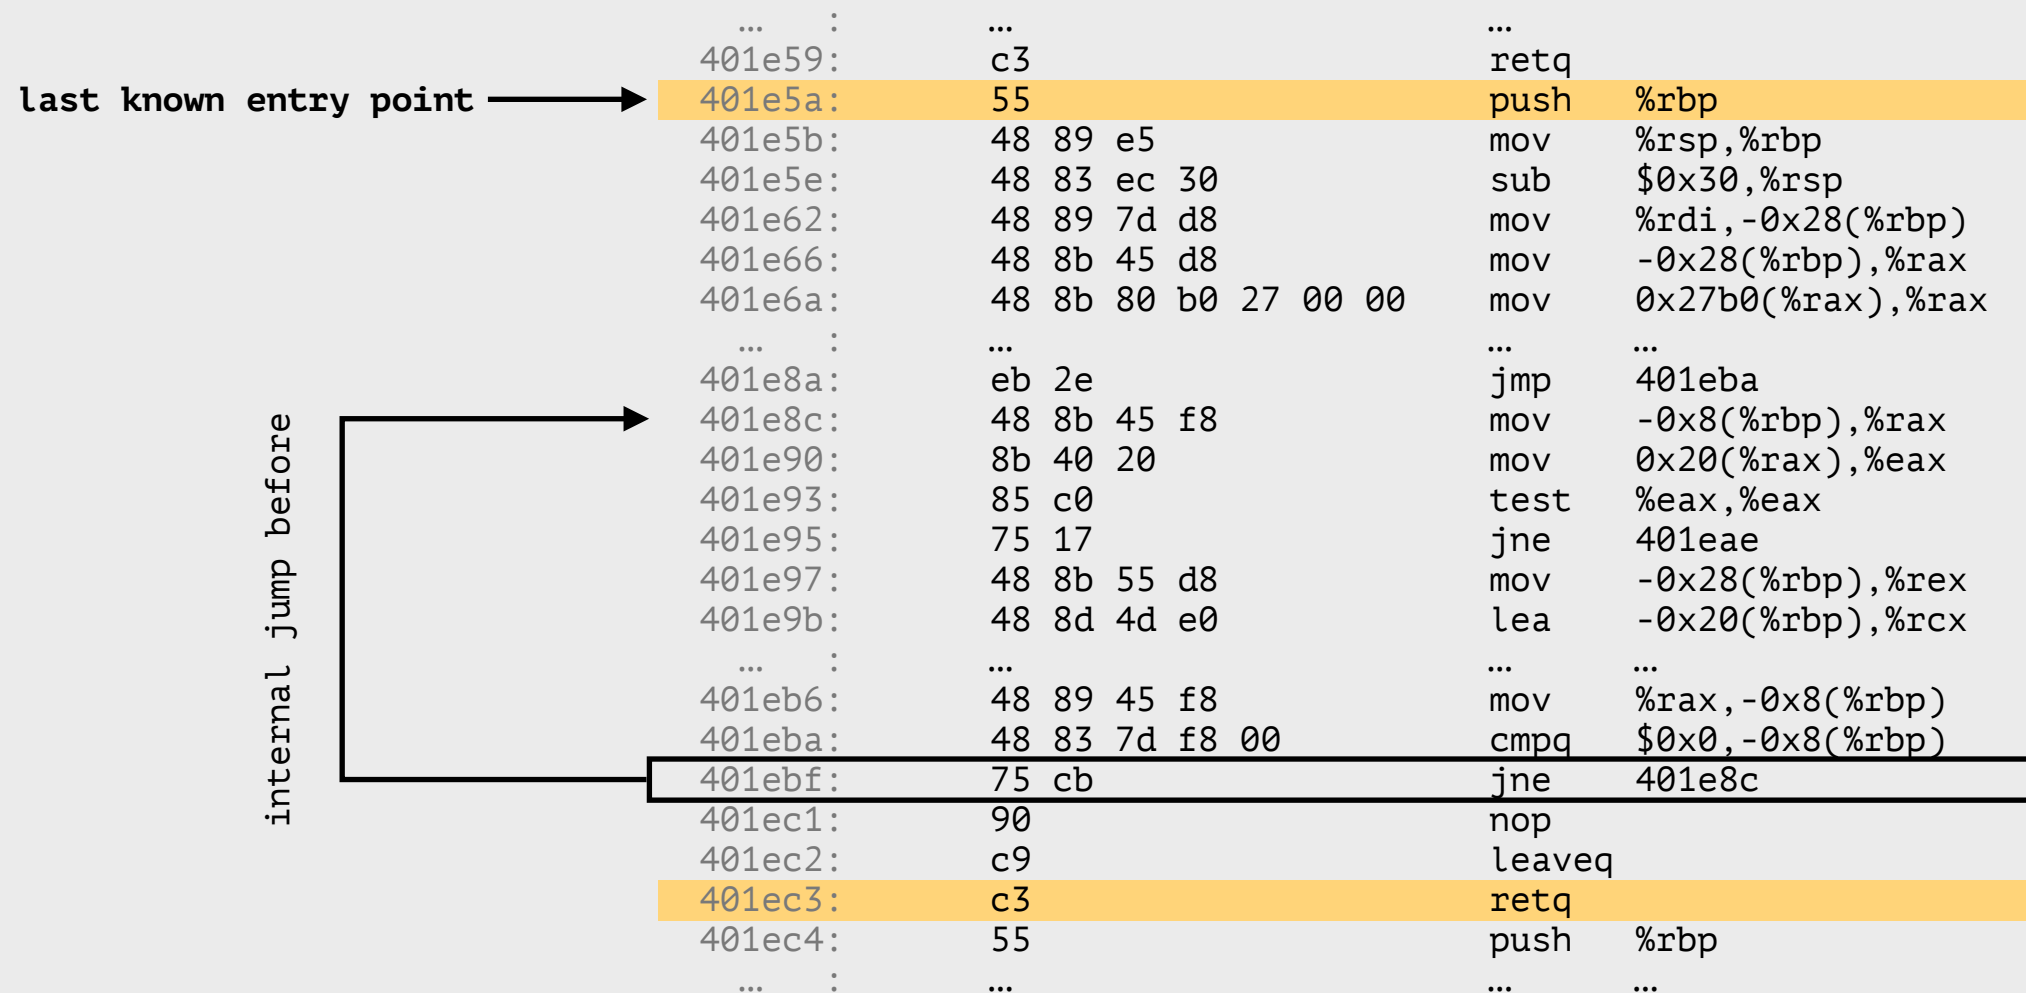
— precision over time

HEURISTICS

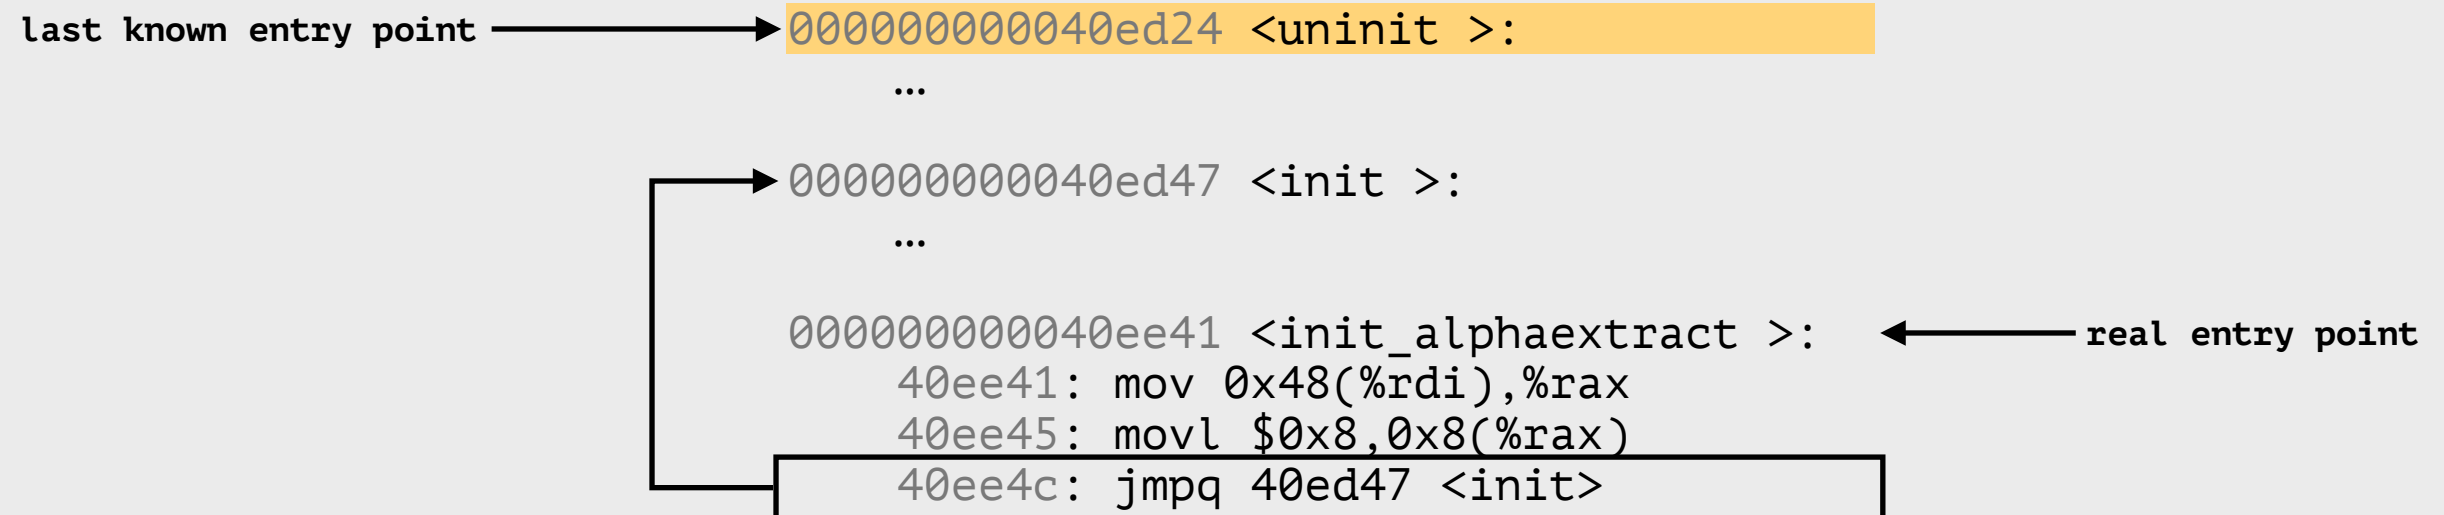


INTERNAL JUMP BEFORE

A jump to a location that is between the current entry point and the current %rip **is not a call**.

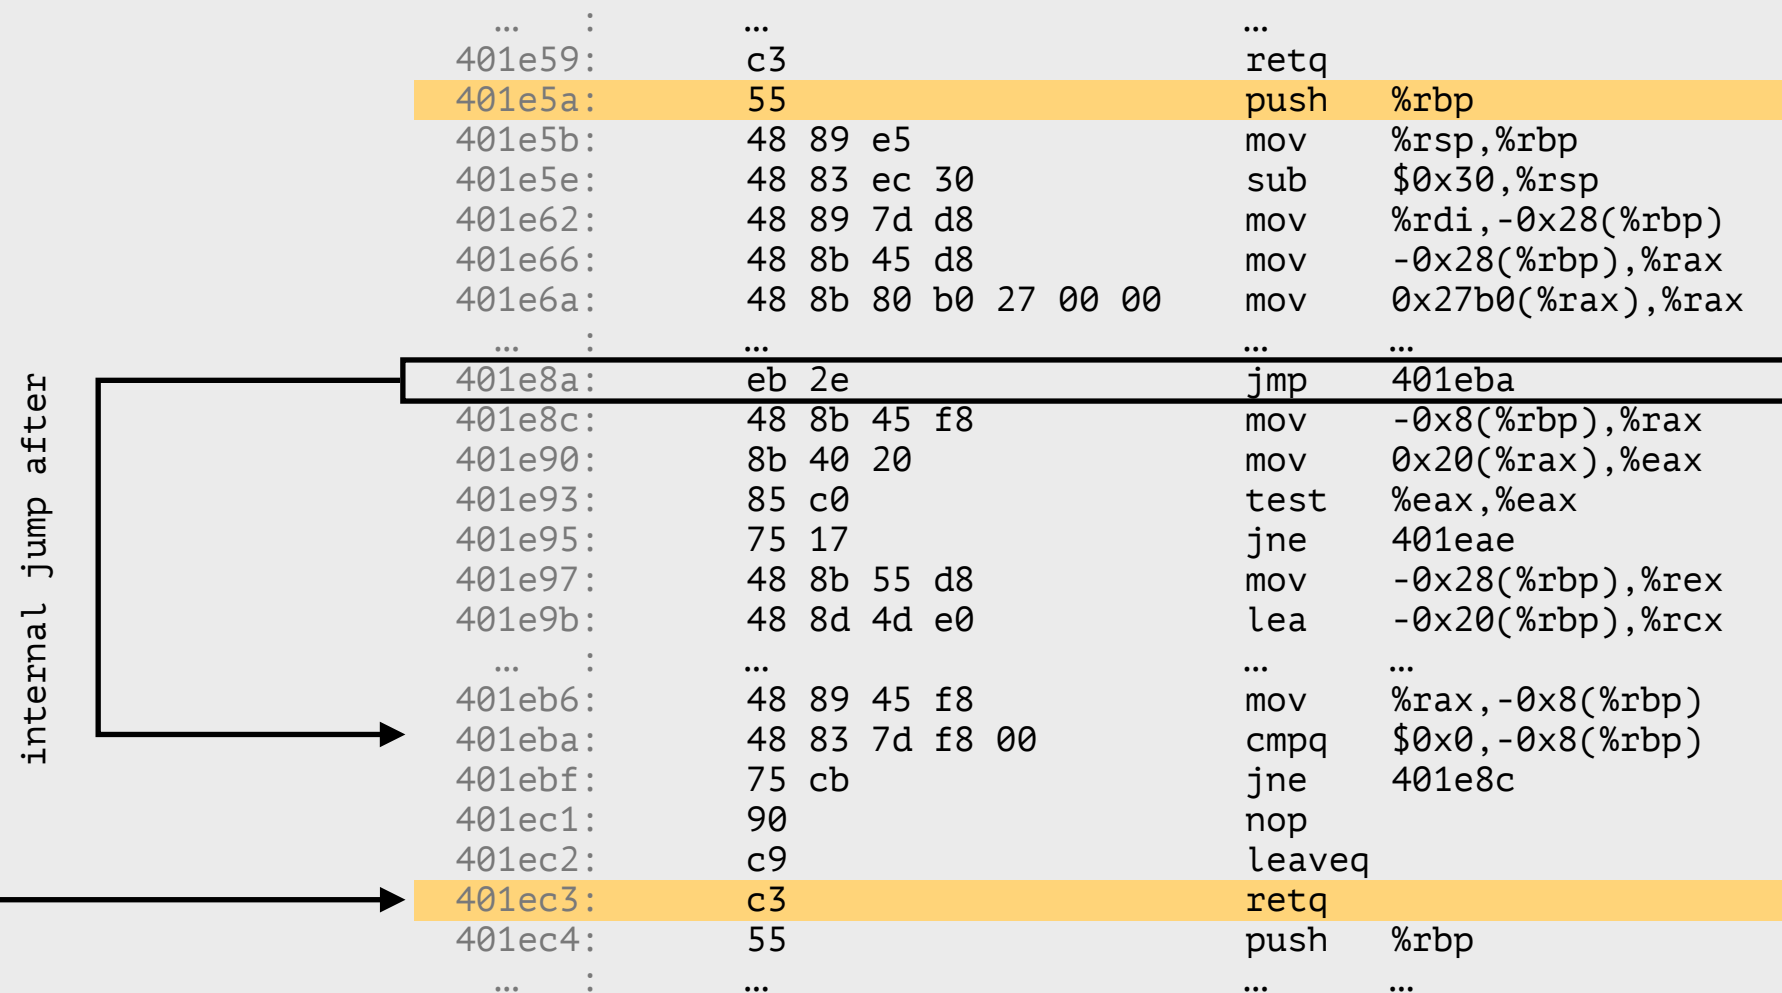


ERROR EXAMPLE



INTERNAL JUMP AFTER

A jump to a location that is between the current %rip and a known ret of the current function **is not a call**.



STACK INCONSISTENCY (IDEA)

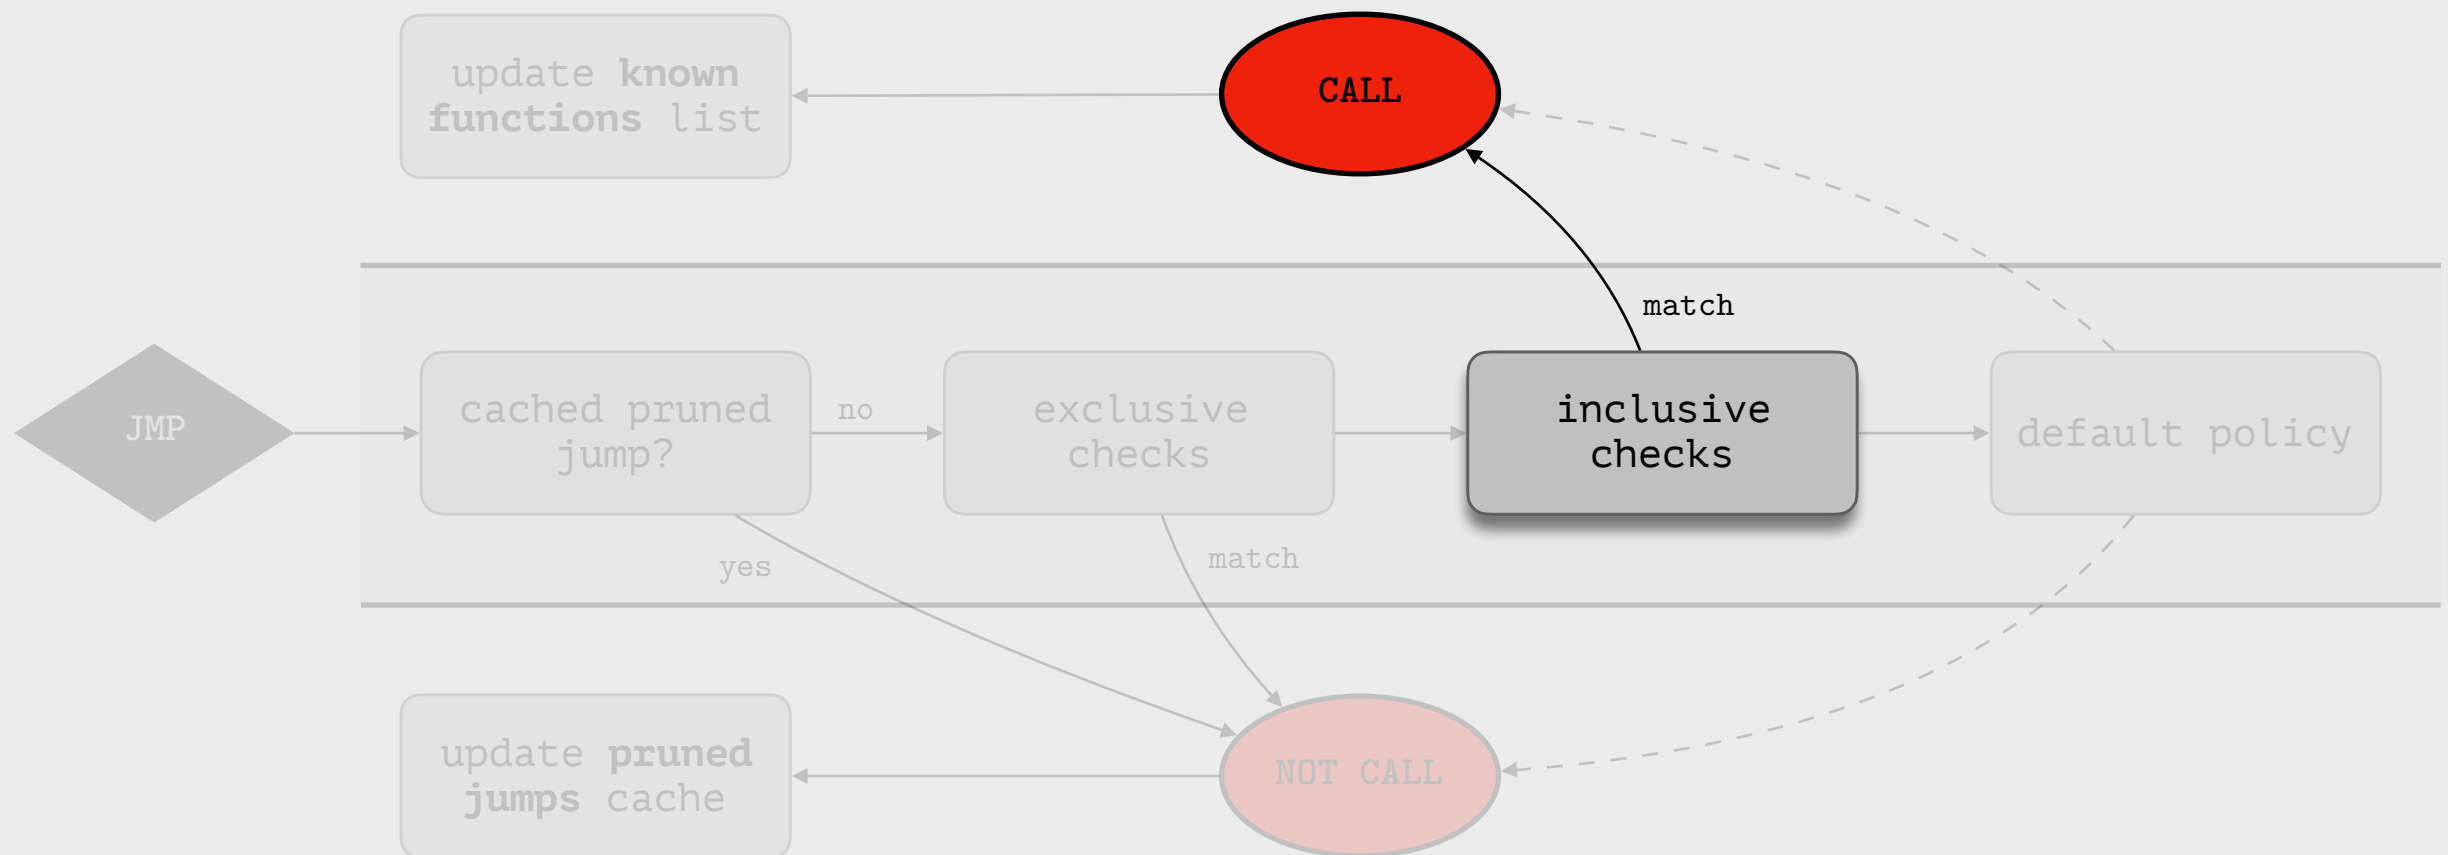
A jump occurring when `%ebp ≠ %esp` is not a (tail) call.

```
push    %rbp
mov     %rsp,%rbp
sub     $0x30,%rsp
push    ...
pop     ...
push    ...
pop     ...
push    ...
jmp     XXX
```

4 push / 2 pop => NOT A CALL

```
push    %rbp
mov     %rsp,%rbp
sub     $0x30,%rsp
push    ...
push    ...
pop     ...
pop     ...
pop     ...
jmp     XXX
```

3 push / 3 pop => MAYBE A CALL



KNOWN ENTRY POINT

A jump to a known entry point is a call.

EXTERNAL JUMP BEFORE

A jump to a location that is before the current entry point is a call.

CROSS ENTRY POINT

A jump crossing a known entry point is a call.

EXPERIMENTS

introducing iCi

approach

heuristics

experiments

methodology

METHODOLOGY

ORACLE

INPUT: entry points of functions

ALGORITHM: for each non-sequential update of %rip,
check its new value. If entry point => call.

REMARKS: existing corner cases, but best oracle we have

JCALL

ALGORITHM: CALL instructions + JMP to .plt => calls

REMARKS: low overhead, no false positive

JMP

ALGORITHM: CALL + JMP instructions => calls

REMARKS: high overhead, no false negative

approach

heuristics

experiments

benchmark

BENCHMARK

BINARIES

- coreutils: 98 programs
- binutils: 13 programs
- evince: 1 program
- ffmpeg: 1 program
- SPEC CPU2006 (C/C++): 19 programs

approach

heuristics

experiments

benchmark

OPTIMIZATION LEVEL

for $\{\text{SPEC CPU2006}\}$: -O0 to -O3 with gcc-6.0

COMPILATION

for coreutils: -O0 to -O3 with gcc-6.0 and clang-3.8

approach

heuristics

experiments

results

RESULTS

approach

heuristics

experiments

results

F - SCORE

	-00			-01			-02			-03		
	jcall	jmp	iCi	jcall	jmp	iCi	jcall	jmp	iCi	jcall	jmp	iCi
binutils	1,000	0,223	1,000	1,000	0,222	1,000	0,999	0,222	1,000	1,000	0,221	1,000
evince	0,988	0,831	0,991	0,980	0,796	0,985	0,936	0,803	0,985	0,947	0,818	0,986
coreutils	0,998	0,518	0,999	0,998	0,352	0,998	0,946	0,405	0,998	0,997	0,351	0,998
ffmpeg	0,948	0,227	0,997	0,919	0,183	0,997	0,874	0,182	0,997	0,884	0,181	0,997
SPEC2006 jcall worst	nc	nc	nc	nc	nc	nc	0,848	0,432	1,000	nc	nc	nc
SPEC2006 iCi worst	nc	nc	nc	nc	nc	nc	0,906	0,260	0,947	nc	nc	nc

OVERHEAD

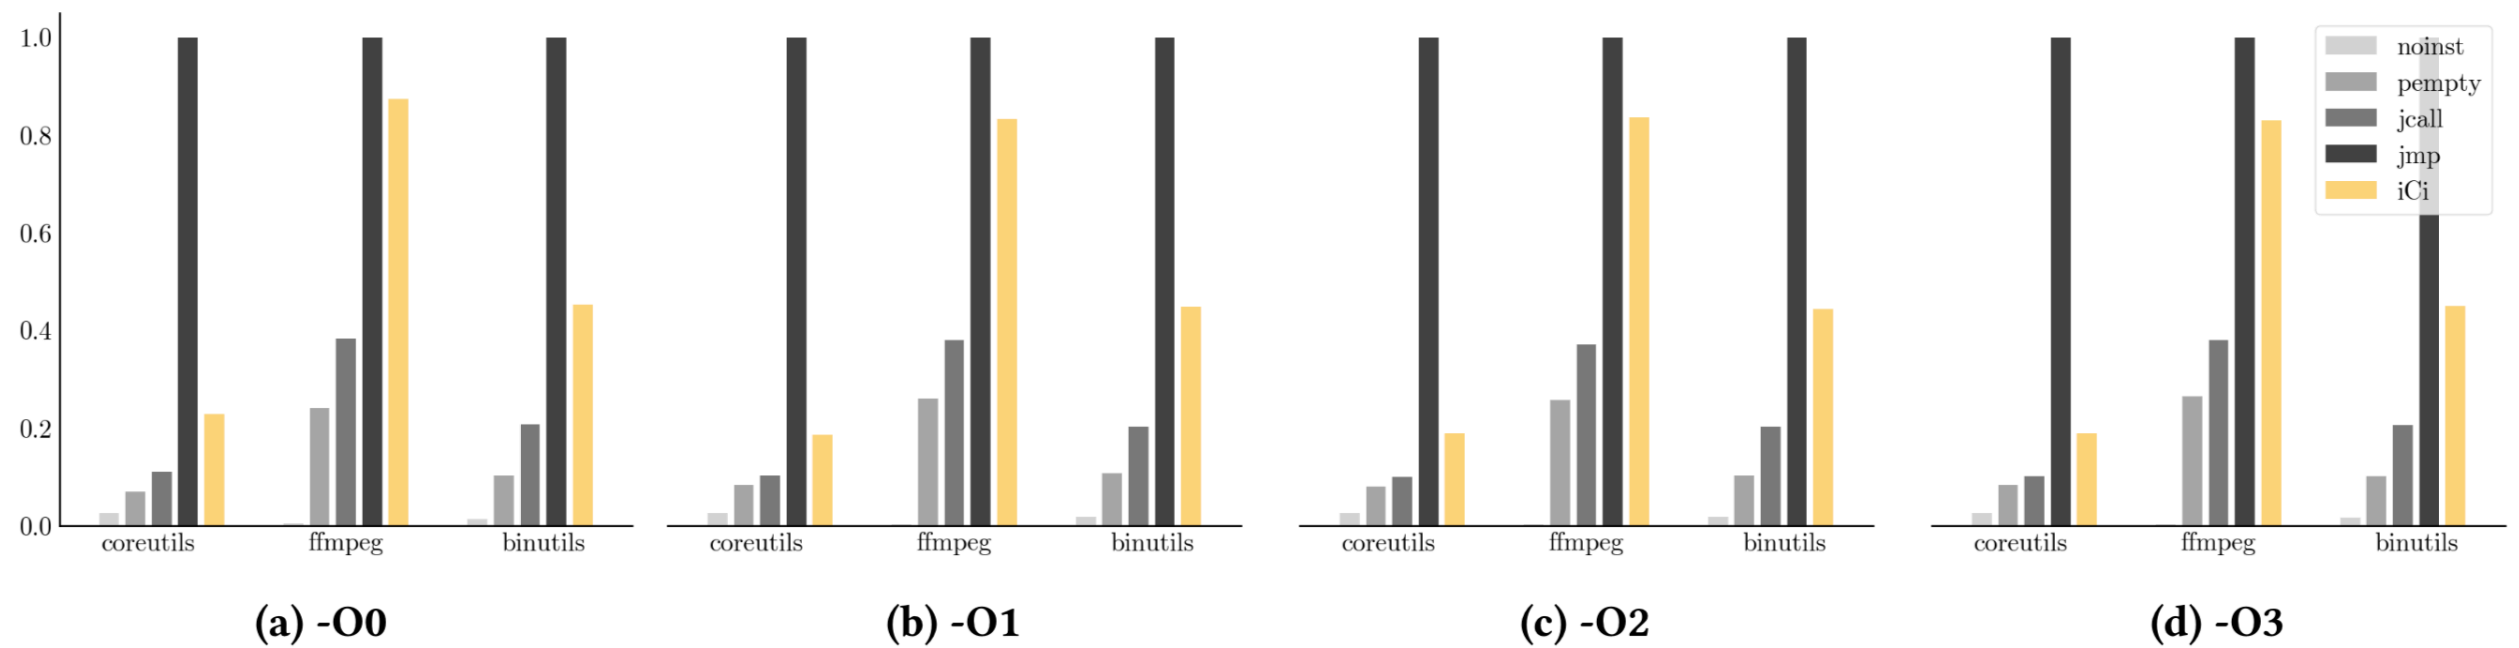


Figure 2: Overhead of each analysis in comparison with jmp for -O2

approach

heuristics

experiments

results

COMPILATION

	-00			-01			-02			-03		
	jcall	jmp	iCi	jcall	jmp	iCi	jcall	jmp	iCi	jcall	jmp	iCi
gcc-6.3	0,998	0,518	0,999	0,998	0,352	0,998	0,946	0,405	0,998	0,997	0,351	0,998
clang-3.8	0,997	0,360	0,998	0,997	0,359	0,998	0,997	0,359	0,998	0,997	0,359	0,998

FINAL WORDS

time to wake everyone up

PROBLEM NOT ADDRESSED BEFORE (TO OUR KNOWLEDGE)

- **TOOLS:** No high-level API to instrument every call (PIN, DynamoRIO, etc.)
- **RESEARCH:** A lot of papers about tail calls... but none about detecting them dynamically

RESULTS INTERPRETATION

iCi is cool! It achieves great f-scores on C/C++ programs, and works with optimization levels and different compilers.

BUT . . .

RESULTS INTERPRETATION

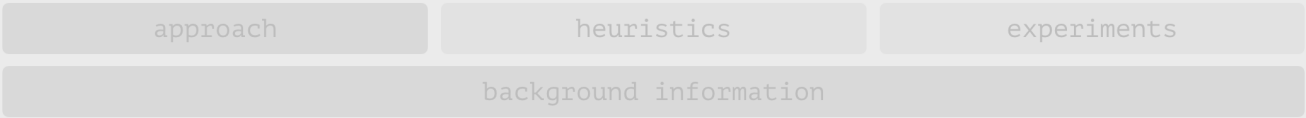
Is it really worth it?

A simpler implementation (jcall) performs almost as good, with a lower overhead.

YES if it is critical to catch all calls

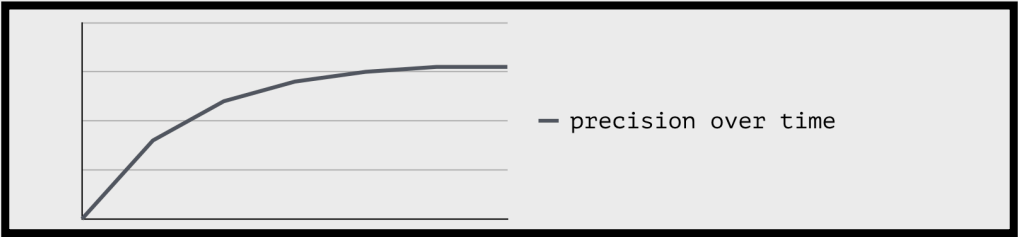
NO otherwise

OPENING



AVAILABLE INFO FOR EACH JMP

information	accessible?
target of the JMP	ok (dynamic)
current %rip, %rsp	ok (dynamic)
current entry point	~ok (depend on the known entry points)
state of stack when last call occurred	~ok (idem)
return locations for the current function	not at first call; then ~ok



17

ACSAC'18

+ overhead decreases over execution

OPENING

BETTER OVER TIME: keep learnt info from an execution to another > should reduce overhead significantly

MIXED ANALYSIS: first static analysis (e.g. Nucleus)

<https://github.com/Frky/iCi>

IMPLEMENTATION

BENCHMARK (EXCEPT SPEC)

EVERYTHING NEEDED TO REPLAY